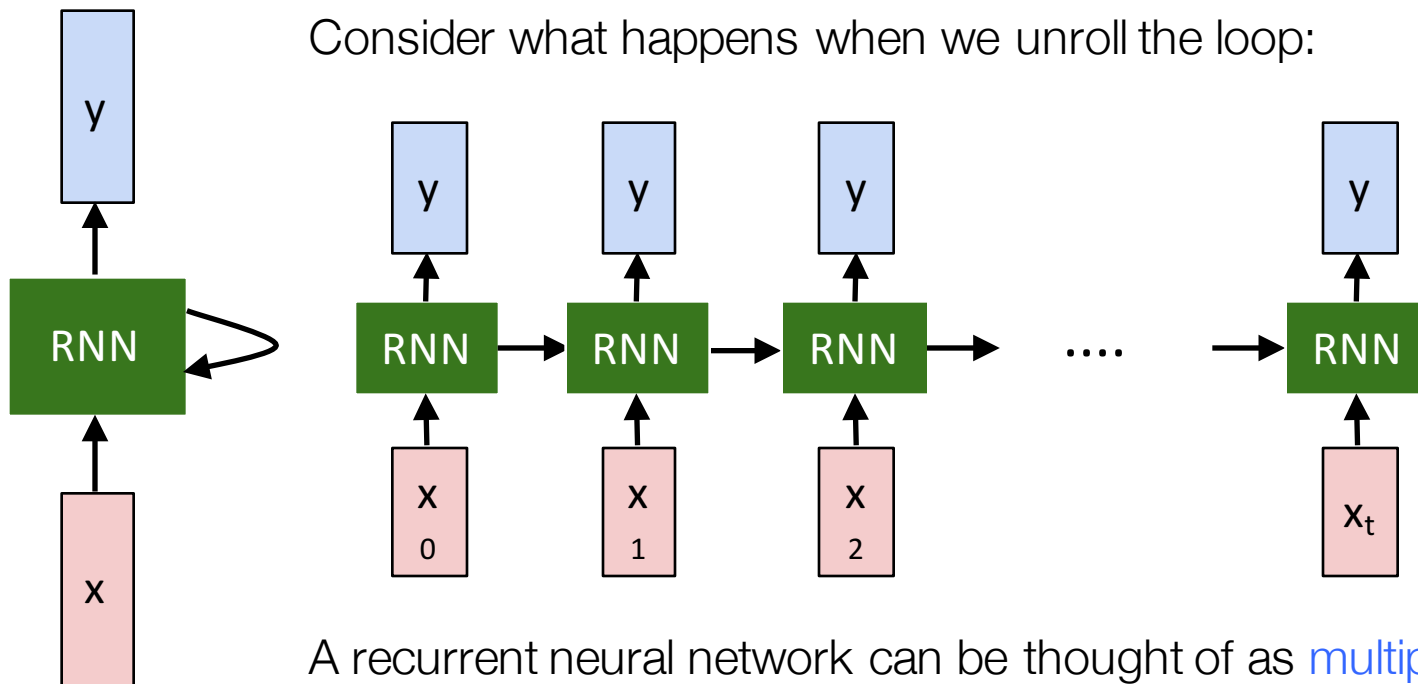


Recurrent Neural Network

Consider what happens when we unroll the loop:



A recurrent neural network can be thought of as **multiple copies of the same network**, each passing a message to a successor.

Recurrent Neural Network

We can process a sequence of vectors x by applying a recurrence formula at every time step:

$$\boxed{h_t} = \boxed{f_W}(\boxed{h_{t-1}}, \boxed{x_t})$$


new state

some function with parameters W

old state


input vector at some time step

Important: the same function and the same set of parameters are used at every time step.



(Vanilla) Recurrent Neural Network

The state consists of a single “*hidden*” vector \mathbf{h} :



$$h_t = f_W(h_{t-1}, x_t)$$




$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

$$y_t = W_{hy}h_t$$



Backpropagation Through Time (BPTT)

- The recurrent model is represented as a multi-layer one (with an unbounded number of layers) and backpropagation is applied on the unrolled model




The problem of long-term dependencies

- (Vanilla) RNNs connect previous information to present task:
- - enough for predicting the next word for “the clouds are in the *sky*”
- - may not be enough when more context is needed
- “I grew up in France... I speak fluent *French*.”




The problem of vanishing gradients

- In a traditional recurrent neural network, during the gradient backpropagation phase, the gradient signal can end up being multiplied a large number of times
- If the gradients are large
 - Exploding gradients, learning diverges
 - **Solution: Clip the gradients to a certain max value.**
- If the gradients are small
 - Vanishing gradients, learning very slow or stops
 - **Solution: introducing memory via LSTM, GRU, etc.**




All recurrent neural networks have the form of a chain of repeating modules of neural network




The repeating module in a standard RNN contains a single layer.

Long Short Term Memory (LSTM)^[Hochreiter & Schmidhuber (1997)]


- A memory cell using logistic and linear units with multiplicative interactions:
- Information gets into the cell whenever its **input** gate is on.
- The information stays in the cell so long as its **forget** gate is on.
- Information can be read from the cell by turning on its **output** gate.




The repeating module in an LSTM contains four interacting layers.



The Core Idea Behind LSTMs : Cell State




Gates are a way to optionally let information through. They are composed out of a sigmoid neural net layer and a pointwise multiplication operation.



An LSTM has three of these gates, to protect and control the cell state.

LSTM : Forget gate



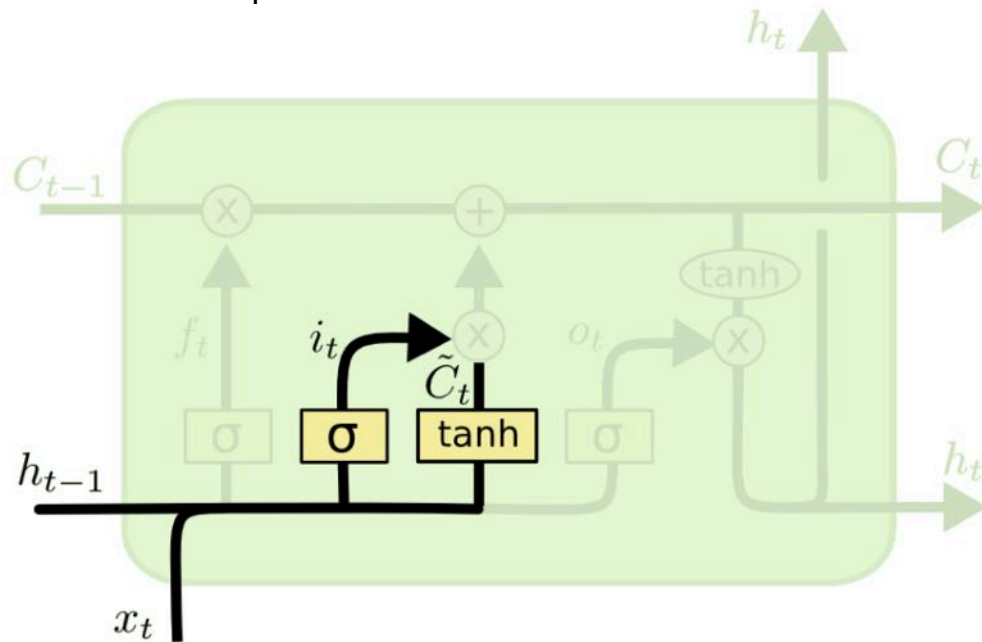
$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$

It looks at h_{t-1} and x_t and outputs a number between 0 and 1 for each number in the cell state C_{t-1} .

A 1 represents **completely keep this** while a 0 represents **completely get rid of this**.

LSTM : Input gate and Cell State

The next step is to decide what new information we're going to store in the cell state.



a sigmoid layer called the **input gate layer** decides which values we'll update.

$$i_t = \sigma (W_i \cdot [h_{t-1}, x_t] + b_i)$$

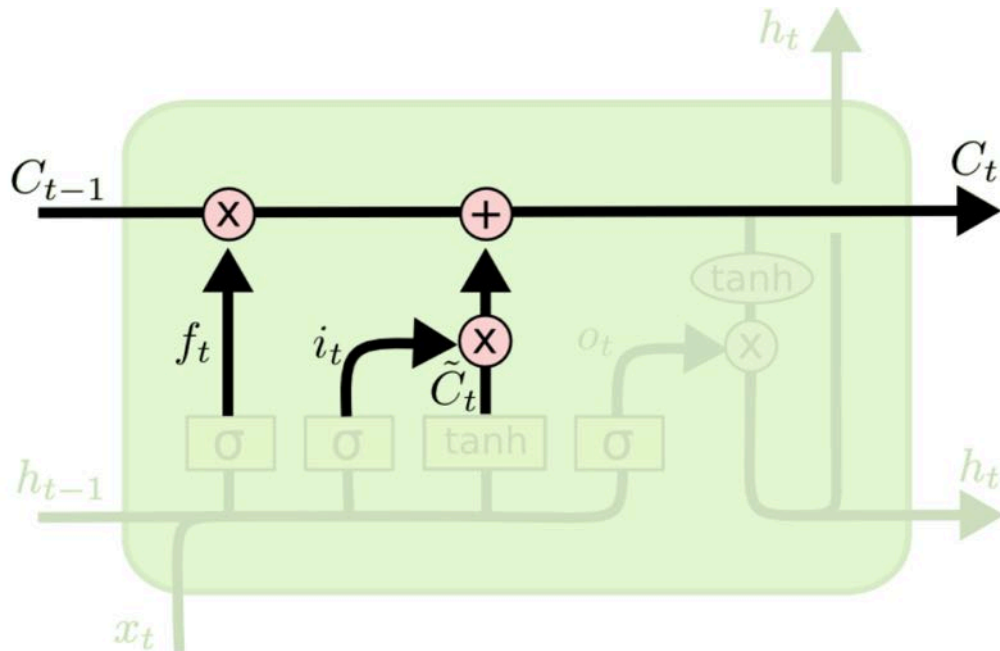
a tanh layer creates a vector of new candidate values, that could be added to the state.

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

LSTM : Input gate and Cell State

It's now time to update the old cell state into the new cell state:

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$




We multiply the old state by f_t forgetting the things we decided to forget earlier.

Then, we add the new candidate values, scaled by how much we decided to update each state value.

LSTM : Output

Finally, we need to decide what we're going to output.



First, we run a sigmoid layer which decides what parts of the cell state we're going to output.

$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$

Then, we put the cell state through tanh (to push the values to be between -1 and 1) and multiply it by the output of the sigmoid gate, so that we only output the parts we decided to.

$$h_t = o_t * \tanh(C_t)$$